

## Assignment 11-1

### 1. Transfer data between MySQL and HDFS (Import and Export) using Sqoop.

#### 1.1 Export data from HDFS to MySQL

Start mysqld service and enter mysql shell. Create new database emp.

```
[acadgild@localhost ~]$ sudo service mysqld start
[sudo] password for acadgild:
Starting mysqld: [ OK ]
[acadgild@localhost ~]$ mysql -uroot
ERROR 1045 (28000): Access denied for user 'root'@'localhost' (using password: NO)
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ mysql -uroot -p
Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 10
Server version: 8.0.3-rc-log MySQL Community Server (GPL)
```

Copyright (c) 2000, 2017, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

```
mysql> show databases;
+-----+
| Database |
+-----+
| information_schema |
| metastore |
| mysql |
| oozie |
| performance_schema |
| simplidb |
| sys |
+-----+
7 rows in set (20.13 sec)

mysql> create database emp;
Query OK, 1 row affected (2.95 sec)

mysql> use emp;
Database changed
```

Check for input file in hdfs

```

[acadgild@localhost ~]$ hadoop fs -ls /
18/02/16 17:03:46 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... us
ing builtin-java classes where applicable
Found 5 items
drwxr-xr-x - acadgild supergroup          0 2018-02-16 16:28 /hbase
drwxr-xr-x - acadgild supergroup          0 2018-02-16 17:01 /sqoop_mysql
drwxr-xr-x - acadgild supergroup          0 2018-02-02 12:49 /sqoopout111
drwxrwx--- - acadgild supergroup          0 2018-02-09 11:35 /tmp
drwxr-xr-x - acadgild supergroup          0 2018-02-09 14:50 /user
[acadgild@localhost ~]$ pwd
/home/acadgild
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -put /home/acadgild/input.txt /sqoop_mysql
18/02/16 17:30:19 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... us
ing builtin-java classes where applicable
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop dfs -cat /sqoop_mysql/in*
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

18/02/16 17:32:11 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... us
ing builtin-java classes where applicable
5,HADOOP,50000,2011-03-21
6,SPARK,60000,2011-03-22
7,JAVA,40000,2011-04-04
8,PIG,50000,2011-01-03
9,HBASE,70000,2011-03-15
10,GIRAF,80000,2011-09-26You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ █

```

Use sqoop command to export from hdfs to mysql

**sqoop export -m 1 --connect jdbc:mysql://localhost/emp --username root --password Root@123 --table acad --export-dir /sqoop\_mysql/input.txt**

```

[acadgild@localhost ~]$ sqoop export -m 1 --connect jdbc:mysql://localhost/emp --username root --password R
oot@123 --table acad --export-dir /sqoop_mysql/input.txt

```

```

18/02/16 21:41:05 INFO impl.YarnClientImpl: Submitted application application_1518776787193_0010
18/02/16 21:41:05 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_151
8776787193_0010/
18/02/16 21:41:05 INFO mapreduce.Job: Running job: job_1518776787193_0010
18/02/16 21:41:21 INFO mapreduce.Job: Job job_1518776787193_0010 running in uber mode : false
18/02/16 21:41:21 INFO mapreduce.Job: map 0% reduce 0%
18/02/16 21:41:33 INFO mapreduce.Job: map 100% reduce 0%
18/02/16 21:41:34 INFO mapreduce.Job: Job job_1518776787193_0010 completed successfully
18/02/16 21:41:34 INFO mapreduce.Job: Counters: 30
    File System Counters
      FILE: Number of bytes read=0
      FILE: Number of bytes written=127522
      FILE: Number of read operations=0
      FILE: Number of large read operations=0
      FILE: Number of write operations=0
      HDFS: Number of bytes read=214
      HDFS: Number of bytes written=0
      HDFS: Number of read operations=4
      HDFS: Number of large read operations=0
      HDFS: Number of write operations=0

```

```

Job Counters
  Launched map tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=8751
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=8751
  Total vcore-milliseconds taken by all map tasks=8751
  Total megabyte-milliseconds taken by all map tasks=8961024
Map-Reduce Framework
  Map input records=6
  Map output records=6
  Input split bytes=123
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=104
  CPU time spent (ms)=2000
  Physical memory (bytes) snapshot=106627072
  Virtual memory (bytes) snapshot=2061332480
  Total committed heap usage (bytes)=32571392
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=0
18/02/16 21:41:34 INFO mapreduce.ExportJobBase: Transferred 214 bytes in 33.887 seconds (6.3151 bytes/sec)
18/02/16 21:41:34 INFO mapreduce.ExportJobBase: Exported 6 records.
You have new mail in /var/spool/mail/acadgild
acadgild@localhost ~]$

```

**Job is successful.**

**Check the export in mysql**

```

[acadgild@localhost ~]$ hadoop fs -cat /sqoop_mysql/input.txt
18/02/16 21:39:34 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform..
ing builtin-java classes where applicable
5,HADOOP,50000
6,SPARK,60000
7,JAVA,40000
8,PIG,50000
9,HBASE,70000
10,GIRAF,80000
[acadgild@localhost ~]$ mysql -uroot -p
Enter password:

Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> select * from acad;
+-----+-----+-----+
| id | name | sal |
+-----+-----+-----+
| 5 | HADOOP | 50000 |
| 6 | SPARK | 60000 |
| 7 | JAVA | 40000 |
| 8 | PIG | 50000 |
| 9 | HBASE | 70000 |
| 10 | GIRAF | 80000 |
+-----+-----+-----+
6 rows in set (0.00 sec)

mysql>

```

## 1.2 Import from MySQL to HDFS

**Select a table to import**

```
mysql> use emp;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A
```

```
Database changed
mysql> show tables
-> ;
```

```
+-----+
| Tables_in_emp |
+-----+
| acad          |
| acad2         |
| empdetails    |
+-----+
3 rows in set (0.00 sec)
```

```
mysql> select * from empdetails;
```

```
+-----+-----+-----+
| id | name   | sal  |
+-----+-----+-----+
| 101 | Aarthi | 15000 |
| 102 | Brindha | 25000 |
| 103 | Hashinee | 50000 |
| 104 | Ashvath | 50000 |
| 105 | Praketh | 40000 |
+-----+-----+-----+
5 rows in set (0.00 sec)
```

```
mysql> █
```

## Grant permission

```
mysql> grant all privileges on *.* to 'root'@'localhost' with grant option;
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> █
```

## Use Sqoop Command to import

```
sqoop import --connect jdbc:mysql://localhost/emp -username root -P --target-dir /sqoopout1 -m 1
```

```
File Edit View Search Terminal Help
[acadgild@localhost ~]$ sqoop import --connect jdbc:mysql://localhost/emp --table empdetails -username root -P^
--target-dir /sqoopout1 -m 1
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin__hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog
jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin__hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo
imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/02/18 21:50:00 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/02/18 21:50:07 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/02/18 21:50:07 INFO tool.CodeGenTool: Beginning code generation
Sun Feb 18 21:50:07 IST 2018 WARN: Establishing SSL connection without server's identity verification is not r
ecommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by
default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServer
Certificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or
set useSSL=true and provide truststore for server certificate verification.
18/02/18 21:50:09 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `empdetails` AS t LIMIT 1
18/02/18 21:50:09 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `empdetails` AS t LIMIT 1
18/02/18 21:50:09 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.
5
Note: /tmp/sqoop-acadgild/compile/115f7c53a89773e7857e417afa09f2f6/empdetails.java uses or overrides a depreca
ted API.
Note: Recompile with -Xlint:deprecation for details.
18/02/18 21:50:14 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/115f7c53a89773e78
57e417afa09f2f6/empdetails.jar
18/02/18 21:50:14 WARN manager.MySQLManager: It looks like you are importing from mysql.
18/02/18 21:50:14 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
18/02/18 21:50:14 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
18/02/18 21:50:14 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
18/02/18 21:50:14 INFO mapreduce.ImportJobBase: Beginning import of empdetails
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log
4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/sl
f4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
```

**Job completed**

```
File Edit View Search Terminal Help
FILE: Number of bytes written=127702
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=87
HDFS: Number of bytes written=100
HDFS: Number of read operations=4
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
Job Counters
  Launched map tasks=1
  Other local map tasks=1
  Total time spent by all maps in occupied slots (ms)=9800
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=9800
  Total vcore-milliseconds taken by all map tasks=9800
  Total megabyte-milliseconds taken by all map tasks=10035200
Map-Reduce Framework
  Map input records=5
  Map output records=5
  Input split bytes=87
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=146
  CPU time spent (ms)=2730
  Physical memory (bytes) snapshot=109985792
  Virtual memory (bytes) snapshot=2063437824
  Total committed heap usage (bytes)=32571392
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=100
18/02/18 21:51:02 INFO mapreduce.ImportJobBase: Transferred 100 bytes in 44.8179 seconds (2.2313 bytes/sec)
18/02/18 21:51:02 INFO mapreduce.ImportJobBase: Retrieved 5 records.
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$
```

## Checking the imported data in HDFS

```
[acadgild@localhost ~]$ hadoop fs -ls /
18/02/18 21:52:35 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using
builtin-java classes where applicable
Found 7 items
drwxr-xr-x - acadgild supergroup 0 2018-02-16 16:28 /hbase
-rw-r--r-- 1 acadgild supergroup 69234933 2018-02-16 23:02 /pigdata
drwxr-xr-x - acadgild supergroup 0 2018-02-16 21:39 /sqoop_mysql
drwxr-xr-x - acadgild supergroup 0 2018-02-18 21:51 /sqoopout1
drwxr-xr-x - acadgild supergroup 0 2018-02-02 12:49 /sqoopout111
drwxrwx--- - acadgild supergroup 0 2018-02-09 11:35 /tmp
drwxr-xr-x - acadgild supergroup 0 2018-02-09 14:50 /user
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -ls /sqoopout1
18/02/18 21:52:53 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using
builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup 0 2018-02-18 21:51 /sqoopout1/ SUCCESS
-rw-r--r-- 1 acadgild supergroup 100 2018-02-18 21:50 /sqoopout1/part-m-00000
[acadgild@localhost ~]$ hadoop fs -cat /sqoopout1/part-m-00000
18/02/18 21:53:24 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using
builtin-java classes where applicable
101,Aarthi,15000.0
102,Brindha,25000.0
103,Hashinee,50000.0
104,Ashvath,50000.0
105,Praketh,40000.0
[acadgild@localhost ~]$
```

## 2.Transfer data between Mysql and Hive (Import and Export only selected columns) using Sqoop.

### 2.1 Import from Mysql to Hive (only selected column)

Select a table which is already present in the mysql

```
mysql> use emp;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A
```

```
Database changed
mysql> show tables;
```

```
+-----+
| Tables_in_emp |
+-----+
| acad          |
| empdetails    |
+-----+
2 rows in set (0.00 sec)
```

```
mysql> describe acad;
```

```
+-----+-----+-----+-----+-----+-----+
| Field | Type      | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+-----+
| id    | int(11)   | YES  |     | NULL    |       |
| name  | varchar(50) | YES  |     | NULL    |       |
| sal   | double    | YES  |     | NULL    |       |
+-----+-----+-----+-----+-----+-----+
3 rows in set (0.28 sec)
```

```
mysql> describe empdetails;
```

```
+-----+-----+-----+-----+-----+-----+
| Field | Type      | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+-----+
| id    | int(11)   | YES  |     | NULL    |       |
| name  | varchar(50) | YES  |     | NULL    |       |
| sal   | double    | YES  |     | NULL    |       |
+-----+-----+-----+-----+-----+-----+
3 rows in set (0.05 sec)
```

```
mysql>
```

```
mysql> select * from empdetails;
```

```
+-----+-----+-----+
| id  | name    | sal  |
+-----+-----+-----+
| 101 | Aarthi  | 15000 |
| 102 | Brindha | 25000 |
| 103 | Hashinee | 50000 |
| 104 | Ashvath | 50000 |
| 105 | Praketh | 40000 |
+-----+-----+-----+
5 rows in set (0.09 sec)
```

**Grant all permission and flush privileges**

```
mysql> grant all on *.* to 'root'@'localhost' with grant option;
Query OK, 0 rows affected (0.24 sec)
```

```
mysql> flush privileges;
Query OK, 0 rows affected (0.40 sec)
```

```
mysql>
```

**use sqoop command**



**sqoop import --connect jdbc:mysql://localhost/emp --username root -P --split-by id --columns id,name --table acad --target-dir /tHive --fields-terminated-by "," --hive-table --hive-table default.acad1**

```
File Edit View Search Terminal Help
[acadgild@localhost ~]$ sqoop import --connect jdbc:mysql://localhost/emp --username root -P --split-by id --columns id,
name --table acad --target-dir /tHive --fields-terminated-by "," --hive-import --create-hive-table --hive-table default.
acad1
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin__hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will
fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin__hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports w
ill fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/02/18 15:55:54 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/02/18 15:55:59 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/02/18 15:55:59 INFO tool.CodeGenTool: Beginning code generation
Sun Feb 18 15:56:00 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended
. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit
option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to
'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore f
or server certificate verification.
18/02/18 15:56:01 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `acad` AS t LIMIT 1
18/02/18 15:56:01 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `acad` AS t LIMIT 1
18/02/18 15:56:01 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadgild/compile/e77f4efc5745298d1ceee01029ec0993/acad.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/02/18 15:56:05 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/e77f4efc5745298d1ceee01029e
c0993/acad.jar
18/02/18 15:56:05 WARN manager.MySQLManager: It looks like you are importing from mysql.
18/02/18 15:56:05 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
18/02/18 15:56:05 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
18/02/18 15:56:05 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
18/02/18 15:56:05 INFO mapreduce.ImportJobBase: Beginning import of acad
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5
.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/S
taticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/02/18 15:56:06 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-j
```

## Job Completed

```
18/02/18 15:57:31 INFO mapreduce.ImportJobBase: Transferred 47 bytes in 83.726 seconds (0.5614 bytes/sec)
18/02/18 15:57:31 INFO mapreduce.ImportJobBase: Retrieved 6 records.
Sun Feb 18 15:57:31 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended
. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit
option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to
'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore f
or server certificate verification.
18/02/18 15:57:31 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `acad` AS t LIMIT 1
18/02/18 15:57:31 INFO hive.HiveImport: Loading uploaded data into Hive
18/02/18 15:57:56 INFO hive.HiveImport: SLF4J: Class path contains multiple SLF4J bindings.
18/02/18 15:57:56 INFO hive.HiveImport: SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-
bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
18/02/18 15:57:56 INFO hive.HiveImport: SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/s
lf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
18/02/18 15:57:56 INFO hive.HiveImport: SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/sha
re/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
18/02/18 15:57:56 INFO hive.HiveImport: SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
18/02/18 15:57:56 INFO hive.HiveImport: SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
18/02/18 15:57:58 INFO hive.HiveImport:
18/02/18 15:57:58 INFO hive.HiveImport: Logging initialized using configuration in jar:file:/home/acadgild/install/hive/
apache-hive-2.3.2-bin/lib/hive-common-2.3.2.jar!/hive-log4j2.properties Async: true
18/02/18 15:58:20 INFO hive.HiveImport: OK
18/02/18 15:58:20 INFO hive.HiveImport: Time taken: 17.971 seconds
18/02/18 15:58:21 INFO hive.HiveImport: Loading data to table default.acad1
18/02/18 15:58:22 INFO hive.HiveImport: OK
18/02/18 15:58:22 INFO hive.HiveImport: Time taken: 2.404 seconds
18/02/18 15:58:22 INFO hive.HiveImport: Hive import complete.
18/02/18 15:58:22 INFO hive.HiveImport: Export directory is contains the _SUCCESS file only, removing the directory.
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$
```



## Check in hive table

```
hive> show tables
> ;
OK
acad1
Time taken: 10.341 seconds, Fetched: 1 row(s)
hive> describe acad1;
OK
id                int
name              string
Time taken: 0.517 seconds, Fetched: 2 row(s)
hive> select * from acad1;
OK
5      HADOOP
6      SPARK
7      JAVA
8      PIG
9      HBASE
10     GIRAF
Time taken: 11.681 seconds, Fetched: 6 row(s)
hive> █
```

## 2.2 Export from Hive to Mysql (only selected column)

### Check the table for export

```
hive> show tables
> ;
OK
acad1
Time taken: 10.341 seconds, Fetched: 1 row(s)
hive> describe acad1;
OK
id                int
name              string
Time taken: 0.517 seconds, Fetched: 2 row(s)
hive> select * from acad1;
OK
5      HADOOP
6      SPARK
7      JAVA
8      PIG
9      HBASE
10     GIRAF
Time taken: 11.681 seconds, Fetched: 6 row(s)
hive> █
```

### Create a table in mysql

```
mysql> create table acad2(id int,name varchar(30));
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that corresponds to your MySQL server version
for the right syntax to use near '' at line 1
mysql> create table acad2(id int,name varchar(30));
Query OK, 0 rows affected (0.09 sec)

mysql> select * from acad2;
Empty set (0.00 sec)

mysql> █
```

### Use sqoop command to export data

```
sqoop export --connect jdbc:mysql://localhost/emp --username root --P --table acad2
--export-dir /user/hive/warehouse/acad1
```

```
acadmild@localhost:~$ sqoop export --connect jdbc:mysql://localhost/emp --username root --P --table acad2 --export-dir
/user/hive/warehouse/acad1
Warning: /home/acadmild/install/sqoop/sqoop-1.4.6.bin__hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will
fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadmild/install/sqoop/sqoop-1.4.6.bin__hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports w
ill fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/02/18 17:21:18 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/02/18 17:21:25 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/02/18 17:21:25 INFO tool.CodeGenTool: Beginning code generation
Sun Feb 18 17:21:25 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended
. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit
option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to
'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore f
or server certificate verification.
18/02/18 17:21:27 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `acad2` AS t LIMIT 1
18/02/18 17:21:27 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `acad2` AS t LIMIT 1
18/02/18 17:21:27 INFO orm.CompilationManager: HADOOP MAPRED HOME is /home/acadmild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadmild/compile/6e0d5b1e96d0bddeblad66f3a8c9d3e2/acad2.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/02/18 17:21:31 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadmild/compile/6e0d5b1e96d0bddeblad66f3a8c
9d3e2/acad2.jar
18/02/18 17:21:31 INFO mapreduce.ExportJobBase: Beginning export of acad2
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadmild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5
.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadmild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/S
taticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/02/18 17:21:31 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-j
va classes where applicable
18/02/18 17:21:32 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
18/02/18 17:21:34 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use
mapreduce.reduce.speculative
```

## Job Completed

```
Job Counters
  Launched map tasks=3
  Data-local map tasks=3
  Total time spent by all maps in occupied slots (ms)=96508
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=96508
  Total vcore-milliseconds taken by all map tasks=96508
  Total megabyte-milliseconds taken by all map tasks=98824192
Map-Reduce Framework
  Map input records=6
  Map output records=6
  Input split bytes=654
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=1092
  CPU time spent (ms)=7050
  Physical memory (bytes) snapshot=318554112
  Virtual memory (bytes) snapshot=6188601344
  Total committed heap usage (bytes)=97714176
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=0
18/02/18 17:22:31 INFO mapreduce.ExportJobBase: Transferred 737 bytes in 57.1592 seconds (12.8938 bytes/sec)
18/02/18 17:22:31 INFO mapreduce.ExportJobBase: Exported 6 records.
You have new mail in /var/spool/mail/acadmild
acadmild@localhost ~$
```

## Check in Table

```
mysql> use emp;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A
```

```
Database changed
mysql> show tables;
```

```
+-----+
| Tables_in_emp |
+-----+
| acad          |
| acad2         |
| empdetails    |
+-----+
3 rows in set (0.00 sec)
```

```
mysql> select * from acad2;
```

```
+-----+-----+
| id  | name |
+-----+-----+
| 9   | HBASE |
| 7   | JAVA |
| 5   | HADOOP |
| 8   | PIG |
| 10  | GIRAF |
| 6   | SPARK |
+-----+-----+
6 rows in set (0.00 sec)
```

```
mysql> █
```